

Kaggle_2

March 28, 2025

0.0.1 Dataset Downloading

```
[1]: !pip install wldhx.yadisk-direct  
!curl -L $(yadisk-direct https://disk.yandex.ru/d/-4g-144mqslQXg) -o data.zip  
!unzip data.zip >> /dev/null
```

```
Requirement already satisfied: wldhx.yadisk-direct in  
/usr/local/lib/python3.11/dist-packages (0.0.6)  
Requirement already satisfied: requests in /usr/local/lib/python3.11/dist-  
packages (from wldhx.yadisk-direct) (2.32.3)  
Requirement already satisfied: charset-normalizer<4,>=2 in  
/usr/local/lib/python3.11/dist-packages (from requests->wldhx.yadisk-direct)  
(3.4.1)  
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-  
packages (from requests->wldhx.yadisk-direct) (3.10)  
Requirement already satisfied: urllib3<3,>=1.21.1 in  
/usr/local/lib/python3.11/dist-packages (from requests->wldhx.yadisk-direct)  
(2.3.0)  
Requirement already satisfied: certifi>=2017.4.17 in  
/usr/local/lib/python3.11/dist-packages (from requests->wldhx.yadisk-direct)  
(2025.1.31)  
    % Total    % Received % Xferd  Average Speed   Time     Time     Time  Current  
                   Dload  Upload Total   Spent   Left  Speed  
  0      0     0      0     0       0      0 --:--:--  0:00:01 --:--:--     0  
100  194M  100  194M     0      0    9.8M      0  0:00:19  0:00:19 --:--:-- 14.1M  
replace content/kaggle_dataset_3/train/malignant/malignant (127).png? [y]es,  
[n]o, [A]ll, [N]one, [r]ename: A
```

```
[2]: !pip install -q segmentation-models-pytorch --no-cache-dir  
!pip install albumentations  
!pip install git+https://github.com/lucasb-eyer/pydensecrf.git
```

```
Requirement already satisfied: albumentations in /usr/local/lib/python3.11/dist-  
packages (2.0.5)  
Requirement already satisfied: numpy>=1.24.4 in /usr/local/lib/python3.11/dist-  
packages (from albumentations) (1.26.4)  
Requirement already satisfied: scipy>=1.10.0 in /usr/local/lib/python3.11/dist-  
packages (from albumentations) (1.14.1)  
Requirement already satisfied: PyYAML in /usr/local/lib/python3.11/dist-packages
```

```
(from albumentations) (6.0.2)
Requirement already satisfied: pydantic>=2.9.2 in
/usr/local/lib/python3.11/dist-packages (from albumentations) (2.10.6)
Requirement already satisfied: albucore==0.0.23 in
/usr/local/lib/python3.11/dist-packages (from albumentations) (0.0.23)
Requirement already satisfied: opencv-python-headless>=4.9.0.80 in
/usr/local/lib/python3.11/dist-packages (from albumentations) (4.11.0.86)
Requirement already satisfied: stringzilla>=3.10.4 in
/usr/local/lib/python3.11/dist-packages (from albucore==0.0.23->albumentations)
(3.12.3)
Requirement already satisfied: simsimd>=5.9.2 in /usr/local/lib/python3.11/dist-
packages (from albucore==0.0.23->albumentations) (6.2.1)
Requirement already satisfied: annotated-types>=0.6.0 in
/usr/local/lib/python3.11/dist-packages (from pydantic>=2.9.2->albumentations)
(0.7.0)
Requirement already satisfied: pydantic-core==2.27.2 in
/usr/local/lib/python3.11/dist-packages (from pydantic>=2.9.2->albumentations)
(2.27.2)
Requirement already satisfied: typing-extensions>=4.12.2 in
/usr/local/lib/python3.11/dist-packages (from pydantic>=2.9.2->albumentations)
(4.12.2)
Collecting git+https://github.com/lucasb-eyer/pydensecrf.git
  Cloning https://github.com/lucasb-eyer/pydensecrf.git to /tmp/pip-req-build-
ihw8vu2w
    Running command git clone --filter=blob:none --quiet
https://github.com/lucasb-eyer/pydensecrf.git /tmp/pip-req-build-ihw8vu2w
      Resolved https://github.com/lucasb-eyer/pydensecrf.git to commit
2723c7fa4f2ead16ae1ce3d8afe977724bb8f87f
    Installing build dependencies ... done
    Getting requirements to build wheel ... done
    Preparing metadata (pyproject.toml) ... done
```

```
[ ]: !pip install wandb
```

```
[ ]: import wandb
wandb.login()
```

```
[14]: import albumentations as A
from albumentations.pytorch import ToTensorV2
import os
import glob
import cv2
import numpy as np
import torch
from torch.utils.data import Dataset
import matplotlib.pyplot as plt
```

```

def get_augmentation_pipeline(image_size=(256, 256), augment=True):
    if augment:
        return A.Compose([
            A.HorizontalFlip(p=0.5),
            A.VerticalFlip(p=0.5),
            A.Rotate(limit=30, p=0.5),
            A.ShiftScaleRotate(shift_limit=0.2, scale_limit=0.2,
                                rotate_limit=30, p=0.5),
            A.ElasticTransform(alpha=1, sigma=50, alpha_affine=50, p=0.5),
            A.GridDistortion(p=0.5),
            A.OpticalDistortion(distort_limit=0.2, shift_limit=0.2, p=0.3),
            A.RandomBrightnessContrast(brightness_limit=0.3, contrast_limit=0.
                3, p=0.5),
            A.GaussianBlur(blur_limit=(5, 9), p=0.3),
            A.RandomGamma(gamma_limit=(70, 130), p=0.3),
            A.CLAHE(clip_limit=2.0, tile_grid_size=(8, 8), p=0.3),
            A.Normalize(mean=[0.5], std=[0.5]),
            ToTensorV2(),
        ])
    else:
        return A.Compose([
            A.Normalize(mean=[0.5], std=[0.5]),
            ToTensorV2(),
        ])

class MultiClassBUSIDataset(Dataset):
    def __init__(self, root_dir, image_size=(256, 256), transform=None,
                 samples=None,
                 include_normal=True, augment=False):
        self.root_dir = root_dir
        self.image_size = image_size
        self.transform = transform
        self.include_normal = include_normal
        self.augment = augment
        self.class_to_idx = {'background': 0, 'benign': 1, 'malignant': 2}
        if samples is not None:
            self.samples = samples
        else:
            self.samples = self._scan_dataset()
        self.augmentation_pipeline = get_augmentation_pipeline(image_size,
                    augment=augment)

    def _scan_dataset(self):
        samples = []
        subfolders = ["benign", "malignant", "normal"] if self.include_normal
        else ["benign", "malignant"]
        for sf in subfolders:

```

```

        folder_path = os.path.join(self.root_dir, sf)
        mask_paths = glob.glob(os.path.join(folder_path, "*_mask.png"))
        for mp in mask_paths:
            img_path = mp.replace("_mask", "")
            if os.path.exists(img_path):
                class_id = 0 if sf == "normal" else 1 if sf == "benign" ↵
        ↵else 2
                samples.append((img_path, mp, class_id))
        return samples

    def __len__(self):
        return len(self.samples)

    def __getitem__(self, idx):
        img_path, mask_path, class_id = self.samples[idx]
        image = cv2.imread(img_path, cv2.IMREAD_GRAYSCALE)
        mask = cv2.imread(mask_path, cv2.IMREAD_GRAYSCALE)
        image = cv2.resize(image, self.image_size, interpolation=cv2.INTER_AREA)
        mask = cv2.resize(mask, self.image_size, interpolation=cv2. ↵
        ↵INTER_NEAREST)
        binary_mask = (mask > 127).astype(np.uint8)
        multi_class_mask = np.zeros_like(binary_mask) if class_id == 0 else ↵
        ↵binary_mask * class_id
        if self.augment and self.augmentation_pipeline is not None:
            augmented = self.augmentation_pipeline(image=image, ↵
            ↵mask=multi_class_mask)
            image = augmented['image']
            mask = augmented['mask']
        else:
            image = torch.from_numpy(image).unsqueeze(0).float() / 255.0
            mask = torch.from_numpy(multi_class_mask).long()
        return image, mask

def create_train_val_datasets(root_dir, val_ratio=0.2, image_size=(256, 256), ↵
    ↵include_normal=True, stratify=True, augment_train=True):
    import os
    import glob
    import random
    from collections import defaultdict
    subfolders = ["benign", "malignant", "normal"] if include_normal else ↵
    ↵["benign", "malignant"]
    all_samples = []
    samples_by_class = defaultdict(list)
    for sf in subfolders:
        folder_path = os.path.join(root_dir, sf)
        mask_paths = glob.glob(os.path.join(folder_path, "*_mask.png"))

```

```

for mp in mask_paths:
    img_path = mp.replace("_mask", "")
    if os.path.exists(img_path):
        class_id = 0 if sf == "normal" else 1 if sf == "benign" else 2
        sample = (img_path, mp, class_id)
        all_samples.append(sample)
        samples_by_class[class_id].append(sample)
train_samples = []
val_samples = []
if stratify:
    for class_id, samples in samples_by_class.items():
        random.shuffle(samples)
        val_count = int(len(samples) * val_ratio)
        val_samples.extend(samples[:val_count])
        train_samples.extend(samples[val_count:])
else:
    random.shuffle(all_samples)
    val_count = int(len(all_samples) * val_ratio)
    val_samples = all_samples[:val_count]
    train_samples = all_samples[val_count:]
train_dataset = MultiClassBUSIDataset(root_dir, image_size=image_size, ▾
samples=train_samples, include_normal=include_normal, augment=augment_train)
val_dataset = MultiClassBUSIDataset(root_dir, image_size=image_size, ▾
samples=val_samples, include_normal=include_normal, augment=False)
return train_dataset, val_dataset

root_dir='/content/content/kaggle_dataset_3/train'
train_dataset, val_dataset = create_train_val_datasets(root_dir=root_dir, ▾
val_ratio=0.2, image_size=(256, 256), include_normal=True, stratify=True, ▾
augment_train=True)
image, mask = train_dataset[0]
import matplotlib.pyplot as plt
plt.figure(figsize=(10, 5))
plt.subplot(1, 2, 1)
plt.imshow(image.squeeze(), cmap='gray')
plt.title("Augmented Image")
plt.axis('off')
plt.subplot(1, 2, 2)
plt.imshow(mask, cmap='gray')
plt.title("Augmented Mask")
plt.axis('off')
plt.show()

```


/usr/local/lib/python3.11/dist-packages/albumentations/core/validation.py:87:
UserWarning: ShiftScaleRotate is a special case of Affine transform. Please use
Affine transform instead.

 original_init(self, **validated_kwargs)

```

<ipython-input-14-4b6a996d9396>:18: UserWarning: Argument(s) 'alpha_affine' are
not valid for transform ElasticTransform
    A.ElasticTransform(alpha=1, sigma=50, alpha_affine=50, p=0.5),
<ipython-input-14-4b6a996d9396>:20: UserWarning: Argument(s) 'shift_limit' are
not valid for transform OpticalDistortion
    A.OpticalDistortion(distort_limit=0.2, shift_limit=0.2, p=0.3),

```



```

[19]: import os
import glob
import cv2
import numpy as np
import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.utils.data import Dataset, DataLoader
from torch.optim.lr_scheduler import ReduceLROnPlateau, OneCycleLR
from torch.utils.tensorboard import SummaryWriter
import time
import matplotlib.pyplot as plt
import segmentation_models_pytorch as smp
import albumentations as A
from albumentations.pytorch import ToTensorV2
import pydensecrf.densecrf as dcrf
from pydensecrf.utils import unary_from_softmax

# Set device
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(f"Using device: {device}")

```

```

# Define the enhanced U-Net model with a pre-trained encoder
def get_enhanced_unet(encoder_name='resnet50', encoder_weights='imagenet',  

    ↪in_channels=1, out_channels=3):
    model = smp.Unet(  

        encoder_name=encoder_name,  

        encoder_weights=encoder_weights,  

        in_channels=in_channels,  

        classes=out_channels,  

        decoder_attention_type='scse',  

        activation='softmax2d',  

    )
    return model

# Initialize the enhanced model
model = get_enhanced_unet(encoder_name='resnet50', encoder_weights='imagenet',  

    ↪in_channels=1, out_channels=3)
model = model.to(device)
print(model)

# Define the combined Dice and Focal Loss with class weights
class DiceFocalLoss(nn.Module):
    def __init__(self, alpha=0.5, gamma=2.0, smooth=1e-6, class_weights=None):
        super(DiceFocalLoss, self).__init__()
        self.alpha = alpha
        self.gamma = gamma
        self.smooth = smooth
        self.class_weights = class_weights

    def forward(self, inputs, targets):
        targets = targets.long()
        inputs_softmax = F.softmax(inputs, dim=1)
        targets_one_hot = F.one_hot(targets, num_classes=inputs.shape[1]).  

            ↪permute(0, 3, 1, 2).float()
        intersection = (inputs_softmax * targets_one_hot).sum(dim=(2, 3))
        union = inputs_softmax.sum(dim=(2, 3)) + targets_one_hot.sum(dim=(2, 3))
        dice_loss = 1 - (2 * intersection + self.smooth) / (union + self.smooth)
        ce_loss = F.cross_entropy(inputs, targets, reduction='none',  

            ↪weight=self.class_weights)
        pt = torch.exp(-ce_loss)
        focal_loss = (1 - pt) ** self.gamma * ce_loss
        loss = self.alpha * dice_loss.mean() + (1 - self.alpha) * focal_loss.  

            ↪mean()
        return loss

# Calculate class weights based on dataset
def calculate_class_weights(dataset):
    class_counts = torch.zeros(3)

```

```

for _, mask in dataset:
    class_counts += torch.bincount(mask.flatten(), minlength=3)
class_weights = 1.0 / (class_counts + 1e-6)
class_weights /= class_weights.sum()
return class_weights.to(device)

# Initialize the loss function with class weights
class_weights = calculate_class_weights(train_dataset)
print(f"Class Weights: {class_weights}")
criterion = DiceFocalLoss(alpha=0.5, gamma=2.0, class_weights=class_weights)

# Define the Dice score calculation for a specific class
def calculate_dice_score_for_class(preds, targets, class_id, smooth=1e-6):
    pred_class = (preds == class_id).float()
    target_class = (targets == class_id).float()
    intersection = (pred_class * target_class).sum(dim=(1, 2))
    union = pred_class.sum(dim=(1, 2)) + target_class.sum(dim=(1, 2))
    dice = (2 * intersection + smooth) / (union + smooth)
    return dice.mean()

# Define the Dice score calculation for overall segmentation
def calculate_dice_score(preds, targets, smooth=1e-6):
    preds = preds.float()
    targets = targets.float()
    intersection = (preds * targets).sum(dim=(1, 2))
    union = preds.sum(dim=(1, 2)) + targets.sum(dim=(1, 2))
    dice = (2 * intersection + smooth) / (union + smooth)
    return dice.mean()

# Define the training loop with learning rate warm-up and gradient clipping
def train_enhanced_unet(model, train_loader, val_loader, num_epochs=50,
                        learning_rate=1e-4, device=None):
    import wandb
    wandb.init(project="busi-multiclass-segmentation", config={
        "epochs": num_epochs,
        "batch_size": train_loader.batch_size,
        "learning_rate": learning_rate,
        "architecture": "Enhanced U-Net (ResNet50)",
        "optimizer": "AdamW"
    })
    model, train_loader, val_loader, num_epochs=50, learning_rate=1e-4,
    device=None):
        if device is None:
            device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

    model = model.to(device)

```

```

optimizer = torch.optim.AdamW(model.parameters(), lr=learning_rate, weight_decay=1e-5)
scheduler = OneCycleLR(optimizer, max_lr=learning_rate, steps_per_epoch=len(train_loader), epochs=num_epochs)
writer = SummaryWriter(log_dir=f'runs/unet_{time.strftime("%Y%m%d_%H%M%S")}')
best_val_loss = float('inf')
early_stopping_patience = 10
patience_counter = 0

for epoch in range(num_epochs):
    model.train()
    train_loss = 0.0
    for batch_idx, (data, target) in enumerate(train_loader):
        data, target = data.to(device), target.to(device)
        optimizer.zero_grad()
        output = model(data)
        loss = criterion(output, target)
        loss.backward()
        torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm=1.0)
        optimizer.step()
        scheduler.step()
        train_loss += loss.item()
        writer.add_scalar('Train/Batch Loss', loss.item(), epoch * len(train_loader) + batch_idx)
    train_loss /= len(train_loader)

    model.eval()
    val_loss = 0.0
    val_dice = 0.0
    val_dice_bg = 0.0
    val_dice_benign = 0.0
    val_dice_malignant = 0.0

    with torch.no_grad():
        for data, target in val_loader:
            data, target = data.to(device), target.to(device)
            output = model(data)
            loss = criterion(output, target)
            val_loss += loss.item()
            preds = torch.argmax(output, dim=1)
            dice_bg = calculate_dice_score_for_class(preds, target, class_id=0)
            dice_benign = calculate_dice_score_for_class(preds, target, class_id=1)

```

```

        dice_malignant = calculate_dice_score_for_class(preds, target, ↴
class_id=2)
        val_dice_bg += dice_bg
        val_dice_benign += dice_benign
        val_dice_malignant += dice_malignant
        dice_score = calculate_dice_score(preds, target)
        val_dice += dice_score

    val_loss /= len(val_loader)
    val_dice /= len(val_loader)
    val_dice_bg /= len(val_loader)
    val_dice_benign /= len(val_loader)
    val_dice_malignant /= len(val_loader)

    writer.add_scalar('Train/EPOCH Loss', train_loss, epoch)
    writer.add_scalar('Validation/EPOCH Loss', val_loss, epoch)
    writer.add_scalar('Validation/Dice Score', val_dice, epoch)
    writer.add_scalar('Validation/Dice Background', val_dice_bg, epoch)
    writer.add_scalar('Validation/Dice Benign', val_dice_benign, epoch)
    writer.add_scalar('Validation/Dice Malignant', val_dice_malignant, ↴
epoch)
    ↪

    if val_loss < best_val_loss:
        best_val_loss = val_loss
        patience_counter = 0
        torch.save(model.state_dict(), 'best_model.pth')
    else:
        patience_counter += 1
        if patience_counter >= early_stopping_patience:
            print(f"Early stopping at epoch {epoch}")
            break

    print(f"Epoch {epoch + 1}/{num_epochs}, Train Loss: {train_loss:.4f}, ↴
Val Loss: {val_loss:.4f}")
    print(f"Validation Dice Scores:")
    print(f" - Background: {val_dice_bg:.5f}")
    print(f" - Benign: {val_dice_benign:.5f}")
    print(f" - Malignant: {val_dice_malignant:.5f}")
    print(f" - Overall: {val_dice:.5f}")
    wandb.log({
        "Epoch": epoch,
        "Train Loss": train_loss,
        "Val Loss": val_loss,
        "Val Dice Overall": val_dice,
        "Val Dice Background": val_dice_bg,
        "Val Dice Benign": val_dice_benign,
        "Val Dice Malignant": val_dice_malignant
    })

```

```

    })

writer.close()
return model

# Define the evaluation function with additional metrics
def evaluate_model(model, val_loader, device=None):
    if device is None:
        device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

    model.eval()
    val_loss = 0.0
    val_dice = 0.0
    val_iou = 0.0
    val_precision = 0.0
    val_recall = 0.0

    with torch.no_grad():
        for data, target in val_loader:
            data, target = data.to(device), target.to(device)
            output = model(data)
            loss = criterion(output, target)
            val_loss += loss.item()
            preds = torch.argmax(output, dim=1)
            dice_score = calculate_dice_score(preds, target)
            val_dice += dice_score

            # Calculate IoU, Precision, and Recall
            iou, precision, recall = calculate_metrics(preds, target)
            val_iou += iou
            val_precision += precision
            val_recall += recall

    val_loss /= len(val_loader)
    val_dice /= len(val_loader)
    val_iou /= len(val_loader)
    val_precision /= len(val_loader)
    val_recall /= len(val_loader)

    print(f"Validation Loss: {val_loss:.4f}, Validation Dice: {val_dice:.4f}")
    print(f"Validation IoU: {val_iou:.4f}, Precision: {val_precision:.4f},\nRecall: {val_recall:.4f}")

# Define the function to calculate IoU, Precision, and Recall
def calculate_metrics(preds, targets, num_classes=3, smooth=1e-6):
    iou = 0.0
    precision = 0.0

```

```

recall = 0.0

for class_id in range(num_classes):
    pred_class = (preds == class_id).float()
    target_class = (targets == class_id).float()
    intersection = (pred_class * target_class).sum()
    union = pred_class.sum() + target_class.sum() - intersection
    iou += (intersection + smooth) / (union + smooth)
    precision += (intersection + smooth) / (pred_class.sum() + smooth)
    recall += (intersection + smooth) / (target_class.sum() + smooth)

iou /= num_classes
precision /= num_classes
recall /= num_classes

return iou, precision, recall

# Create train and validation datasets
train_dataset, val_dataset = create_train_val_datasets(
    root_dir='/content/content/kaggle_dataset_3/train',
    val_ratio=0.2,
    image_size=(256, 256),
    include_normal=True,
    stratify=True,
    augment_train=True
)

# Create data loaders
train_loader = DataLoader(train_dataset, batch_size=16, shuffle=True,□
                           num_workers=0)
val_loader = DataLoader(val_dataset, batch_size=16, shuffle=False,□
                           num_workers=0)

# Fix multiprocessing issue in Jupyter
if __name__ == "__main__":
    train_enhanced_unet(model=model, train_loader=train_loader,□
                        val_loader=val_loader, num_epochs=50, learning_rate=1e-4, device=device)
    evaluate_model(model, val_loader, device=device)

```

Using device: cuda

```

Unet(
  (encoder): ResNetEncoder(
    (conv1): Conv2d(1, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3),
bias=False)
    (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)

```

```

(maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1,
ceil_mode=False)
(layer1): Sequential(
(0): Bottleneck(
(conv1): Conv2d(64, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
(bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(relu): ReLU(inplace=True)
(downsample): Sequential(
(0): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
(1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
)
)
(1): Bottleneck(
(conv1): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
(bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(relu): ReLU(inplace=True)
)
(2): Bottleneck(
(conv1): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
(bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(relu): ReLU(inplace=True)
)
)
)

```

```

(layer2): Sequential(
    (0): Bottleneck(
        (conv1): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (downsample): Sequential(
            (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
            (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
    )
    (1): Bottleneck(
        (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
    )
    (2): Bottleneck(
        (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
    )
    (3): Bottleneck(
        (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,

```

```

track_running_stats=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
)
)
(layer3): Sequential(
    (0): Bottleneck(
        (conv1): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (downsample): Sequential(
            (0): Conv2d(512, 1024, kernel_size=(1, 1), stride=(2, 2), bias=False)
            (1): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
    )
    (1): Bottleneck(
        (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
    )
    (2): Bottleneck(
        (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1),

```

```

bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1),
bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
)
(3): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1),
bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1),
bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
)
(4): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1),
bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1),
bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
)
(5): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1),
bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

```

```

        (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
    )
)
(layer4): Sequential(
    (0): Bottleneck(
        (conv1): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (downsample): Sequential(
            (0): Conv2d(1024, 2048, kernel_size=(1, 1), stride=(2, 2), bias=False)
            (1): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
    )
    (1): Bottleneck(
        (conv1): Conv2d(2048, 512, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
    )
)
(2): Bottleneck(

```

```

(conv1): Conv2d(2048, 512, kernel_size=(1, 1), stride=(1, 1),
bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
    )
)
)
)
(decoder): UnetDecoder(
    (center): Identity()
    (blocks): ModuleList(
        (0): DecoderBlock(
            (conv1): Conv2dReLU(
                (0): Conv2d(3072, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
                (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
                (2): ReLU(inplace=True)
            )
            (attention1): Attention(
                (attention): SCSEModule(
                    (cSE): Sequential(
                        (0): AdaptiveAvgPool2d(output_size=1)
                        (1): Conv2d(3072, 192, kernel_size=(1, 1), stride=(1, 1))
                        (2): ReLU(inplace=True)
                        (3): Conv2d(192, 3072, kernel_size=(1, 1), stride=(1, 1))
                        (4): Sigmoid()
                    )
                    (sSE): Sequential(
                        (0): Conv2d(3072, 1, kernel_size=(1, 1), stride=(1, 1))
                        (1): Sigmoid()
                    )
                )
            )
        )
        (conv2): Conv2dReLU(
            (0): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
            (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
            (2): ReLU(inplace=True)

```

```

)
(attention2): Attention(
    (attention): SCSEModule(
        (cSE): Sequential(
            (0): AdaptiveAvgPool2d(output_size=1)
            (1): Conv2d(256, 16, kernel_size=(1, 1), stride=(1, 1))
            (2): ReLU(inplace=True)
            (3): Conv2d(16, 256, kernel_size=(1, 1), stride=(1, 1))
            (4): Sigmoid()
        )
        (sSE): Sequential(
            (0): Conv2d(256, 1, kernel_size=(1, 1), stride=(1, 1))
            (1): Sigmoid()
        )
    )
)
)
)
)
)
)
(1): DecoderBlock(
    (conv1): Conv2dReLU(
        (0): Conv2d(768, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU(inplace=True)
    )
    (attention1): Attention(
        (attention): SCSEModule(
            (cSE): Sequential(
                (0): AdaptiveAvgPool2d(output_size=1)
                (1): Conv2d(768, 48, kernel_size=(1, 1), stride=(1, 1))
                (2): ReLU(inplace=True)
                (3): Conv2d(48, 768, kernel_size=(1, 1), stride=(1, 1))
                (4): Sigmoid()
            )
            (sSE): Sequential(
                (0): Conv2d(768, 1, kernel_size=(1, 1), stride=(1, 1))
                (1): Sigmoid()
            )
        )
    )
    (conv2): Conv2dReLU(
        (0): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU(inplace=True)
    )
    (attention2): Attention(

```



```

        (2): ReLU(inplace=True)
        (3): Conv2d(2, 32, kernel_size=(1, 1), stride=(1, 1))
        (4): Sigmoid()
    )
    (sSE): Sequential(
        (0): Conv2d(32, 1, kernel_size=(1, 1), stride=(1, 1))
        (1): Sigmoid()
    )
)
)
)
)
)
(4): DecoderBlock(
    (conv1): Conv2dReLU(
        (0): Conv2d(32, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
        (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU(inplace=True)
    )
    (attention1): Attention(
        (attention): SCSEModule(
            (cSE): Sequential(
                (0): AdaptiveAvgPool2d(output_size=1)
                (1): Conv2d(32, 2, kernel_size=(1, 1), stride=(1, 1))
                (2): ReLU(inplace=True)
                (3): Conv2d(2, 32, kernel_size=(1, 1), stride=(1, 1))
                (4): Sigmoid()
            )
            (sSE): Sequential(
                (0): Conv2d(32, 1, kernel_size=(1, 1), stride=(1, 1))
                (1): Sigmoid()
            )
        )
    )
    (conv2): Conv2dReLU(
        (0): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
        (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU(inplace=True)
    )
    (attention2): Attention(
        (attention): SCSEModule(
            (cSE): Sequential(
                (0): AdaptiveAvgPool2d(output_size=1)
                (1): Conv2d(16, 1, kernel_size=(1, 1), stride=(1, 1))
                (2): ReLU(inplace=True)
                (3): Conv2d(1, 16, kernel_size=(1, 1), stride=(1, 1))

```


- Overall: 0.21430

Epoch 4/50, Train Loss: 0.3993, Val Loss: 0.3958

Validation Dice Scores:

- Background: 0.76261
- Benign: 0.10454
- Malignant: 0.04772
- Overall: 0.27235

Epoch 5/50, Train Loss: 0.3947, Val Loss: 0.3880

Validation Dice Scores:

- Background: 0.85068
- Benign: 0.15529
- Malignant: 0.08236
- Overall: 0.36536

Epoch 6/50, Train Loss: 0.3897, Val Loss: 0.3820

Validation Dice Scores:

- Background: 0.85335
- Benign: 0.18421
- Malignant: 0.10087
- Overall: 0.38718

Epoch 7/50, Train Loss: 0.3848, Val Loss: 0.3761

Validation Dice Scores:

- Background: 0.91551
- Benign: 0.17359
- Malignant: 0.12912
- Overall: 0.47743

Epoch 8/50, Train Loss: 0.3805, Val Loss: 0.3765

Validation Dice Scores:

- Background: 0.91404
- Benign: 0.20190
- Malignant: 0.03120
- Overall: 0.40701

Epoch 9/50, Train Loss: 0.3777, Val Loss: 0.3714

Validation Dice Scores:

- Background: 0.95444
- Benign: 0.23462
- Malignant: 0.11896
- Overall: 0.56194

Epoch 10/50, Train Loss: 0.3742, Val Loss: 0.3719

Validation Dice Scores:

- Background: 0.95338
- Benign: 0.21589
- Malignant: 0.17719
- Overall: 0.63167

Epoch 11/50, Train Loss: 0.3719, Val Loss: 0.3697

Validation Dice Scores:

- Background: 0.95434
- Benign: 0.34335
- Malignant: 0.18488

- Overall: 0.46523

Epoch 12/50, Train Loss: 0.3692, Val Loss: 0.3691

Validation Dice Scores:

- Background: 0.92842
- Benign: 0.41659
- Malignant: 0.29032
- Overall: 0.58793

Epoch 13/50, Train Loss: 0.3679, Val Loss: 0.3650

Validation Dice Scores:

- Background: 0.94742
- Benign: 0.34795
- Malignant: 0.61216
- Overall: 0.58344

Epoch 14/50, Train Loss: 0.3664, Val Loss: 0.3646

Validation Dice Scores:

- Background: 0.96494
- Benign: 0.62541
- Malignant: 0.57384
- Overall: 0.73988

Epoch 15/50, Train Loss: 0.3643, Val Loss: 0.3655

Validation Dice Scores:

- Background: 0.96899
- Benign: 0.61503
- Malignant: 0.76270
- Overall: 0.72744

Epoch 16/50, Train Loss: 0.3648, Val Loss: 0.3640

Validation Dice Scores:

- Background: 0.96606
- Benign: 0.59448
- Malignant: 0.64905
- Overall: 0.82228

Epoch 17/50, Train Loss: 0.3624, Val Loss: 0.3668

Validation Dice Scores:

- Background: 0.93908
- Benign: 0.32146
- Malignant: 0.75277
- Overall: 0.53057

Epoch 18/50, Train Loss: 0.3619, Val Loss: 0.3633

Validation Dice Scores:

- Background: 0.96412
- Benign: 0.47819
- Malignant: 0.72946
- Overall: 0.67035

Epoch 19/50, Train Loss: 0.3604, Val Loss: 0.3596

Validation Dice Scores:

- Background: 0.96708
- Benign: 0.58391
- Malignant: 0.78348

- Overall: 0.79009

Epoch 20/50, Train Loss: 0.3606, Val Loss: 0.3616

Validation Dice Scores:

- Background: 0.96697
- Benign: 0.66068
- Malignant: 0.78911
- Overall: 0.75323

Epoch 21/50, Train Loss: 0.3590, Val Loss: 0.3608

Validation Dice Scores:

- Background: 0.94321
- Benign: 0.44235
- Malignant: 0.69924
- Overall: 0.63133

Epoch 22/50, Train Loss: 0.3592, Val Loss: 0.3611

Validation Dice Scores:

- Background: 0.94617
- Benign: 0.45715
- Malignant: 0.75783
- Overall: 0.72914

Epoch 23/50, Train Loss: 0.3605, Val Loss: 0.3675

Validation Dice Scores:

- Background: 0.96502
- Benign: 0.62647
- Malignant: 0.69175
- Overall: 0.58332

Epoch 24/50, Train Loss: 0.3602, Val Loss: 0.3649

Validation Dice Scores:

- Background: 0.96466
- Benign: 0.46426
- Malignant: 0.68162
- Overall: 0.66512

Epoch 25/50, Train Loss: 0.3614, Val Loss: 0.3662

Validation Dice Scores:

- Background: 0.95783
- Benign: 0.54769
- Malignant: 0.71217
- Overall: 0.48435

Epoch 26/50, Train Loss: 0.3598, Val Loss: 0.3612

Validation Dice Scores:

- Background: 0.95458
- Benign: 0.67185
- Malignant: 0.65617
- Overall: 0.75232

Epoch 27/50, Train Loss: 0.3593, Val Loss: 0.3668

Validation Dice Scores:

- Background: 0.96030
- Benign: 0.65275
- Malignant: 0.43291

```

- Overall: 0.62670
Epoch 28/50, Train Loss: 0.3607, Val Loss: 0.3630
Validation Dice Scores:
- Background: 0.96654
- Benign: 0.43019
- Malignant: 0.68258
- Overall: 0.72872
Early stopping at epoch 28
Validation Loss: 0.3651, Validation Dice: 0.5589
Validation IoU: 0.4757, Precision: 0.6201, Recall: 0.8345

```

```
[20]: # Evaluate the model
def evaluate_model(model, val_loader, device=None):
    if device is None:
        device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
    model.eval()
    val_loss = 0.0
    val_dice = 0.0
    with torch.no_grad():
        for data, target in val_loader:
            data, target = data.to(device), target.to(device)
            output = model(data)
            loss = criterion(output, target)
            val_loss += loss.item()
            preds = torch.argmax(output, dim=1)
            dice_score = calculate_dice_score(preds, target)
            val_dice += dice_score
    val_loss /= len(val_loader)
    val_dice /= len(val_loader)
    print(f"Validation Loss: {val_loss:.4f}, Validation Dice: {val_dice:.4f}")

evaluate_model(model, val_loader, device=device)

# Visualize predictions
def visualize_predictions(model, val_loader, device, num_samples=3):
    model.eval()
    with torch.no_grad():
        for i, (data, target) in enumerate(val_loader):
            if i >= num_samples:
                break
            data, target = data.to(device), target.to(device)
            output = model(data)
            preds = torch.argmax(output, dim=1)
            image = data[0].cpu().squeeze().numpy()
            target_mask = target[0].cpu().numpy()
            pred_mask = preds[0].cpu().numpy()
            plt.figure(figsize=(15, 5))

```



```


plt.subplot(1, 3, 1)
plt.imshow(image, cmap='gray')
plt.title("Input Image")
plt.axis('off')
plt.subplot(1, 3, 2)
plt.imshow(target_mask, cmap='jet')
plt.title("Ground Truth")
plt.axis('off')
plt.subplot(1, 3, 3)
plt.imshow(pred_mask, cmap='jet')
plt.title("Prediction")
plt.axis('off')
plt.show()

visualize_predictions(model, val_loader, device=device, num_samples=3)

```

Validation Loss: 0.3651, Validation Dice: 0.5589





```
[21]: import os
import torch
import numpy as np
import pandas as pd
import cv2
from torch.utils.data import Dataset, DataLoader
from tqdm import tqdm

# Define the test dataset class
class BUSITestDataset(Dataset):
    def __init__(self, test_dir, image_size=(256, 256)):
        self.test_dir = test_dir
        self.image_size = image_size
        self.image_files = [f for f in os.listdir(test_dir) if f.lower().
        ↪endswith('.png', '.jpg', '.jpeg', '.tif'))]
        self.image_files.sort()

    def __len__(self):
        return len(self.image_files)

    def __getitem__(self, idx):
        image_name = self.image_files[idx]
        image_path = os.path.join(self.test_dir, image_name)
        image_id = os.path.splitext(image_name)[0]
        original_image = cv2.imread(image_path, cv2.IMREAD_GRAYSCALE)
        original_size = original_image.shape
        image = cv2.resize(original_image, self.image_size, interpolation=cv2.
        ↪INTER_AREA)
        image = image.astype(np.float32) / 255.0
        image_tensor = torch.from_numpy(image).unsqueeze(0)
        return {
            'image': image_tensor,
            'image_id': image_id,
```

```

        'original_h': original_size[0],
        'original_w': original_size[1]
    }

def rle_encode_mask(mask):
    if np.sum(mask) == 0:
        return ''
    pixels = mask.flatten()
    runs = np.where(pixels[1:] != pixels[:-1])[0] + 1
    runs = np.concatenate([[0], runs])
    run_lengths = []
    for i in range(len(runs) - 1):
        if pixels[runs[i]] == 1:
            start = runs[i] + 1
            length = runs[i + 1] - runs[i]
            run_lengths.extend([start, length])
    return ' '.join(str(x) for x in run_lengths)

def combined_encode(masks_dict, delimiter="~"):
    if not masks_dict:
        return ""
    encoded_parts = []
    for class_id, mask in masks_dict.items():
        rle = rle_encode_mask(mask)
        if rle:
            encoded_parts.append(f"{class_id}:{rle}")
    return delimiter.join(encoded_parts)

def generate_submission(model, test_dir, output_file, batch_size=4,
    device=None, image_size=(256, 256)):
    if device is None:
        device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
    test_dataset = BUSITestDataset(test_dir, image_size=image_size)
    test_loader = DataLoader(test_dataset, batch_size=batch_size,
    shuffle=False, num_workers=0)
    model.to(device)
    model.eval()
    results = {'ID': [], 'encoded_pixels': []}
    print("Generating predictions...")
    with torch.no_grad():
        for batch in tqdm(test_loader):
            images = batch['image'].to(device)
            image_ids = batch['image_id']
            original_hs = batch['original_h']
            original_ws = batch['original_w']
            outputs = model(images)
            batch_size = images.shape[0]

```

```

is_multiclass = outputs.shape[1] > 1
if is_multiclass:
    predictions = torch.argmax(outputs, dim=1).cpu().numpy()
    for i in range(batch_size):
        image_id = image_ids[i]
        pred = predictions[i]
        orig_h, orig_w = original_hs[i].item(), original_ws[i].item()
        pred_resized = cv2.resize(pred.astype(np.float32), (orig_w, orig_h), interpolation=cv2.INTER_NEAREST).astype(np.int32)
        masks_dict = {}
        for class_id in range(1, outputs.shape[1]):
            binary_mask = (pred_resized == class_id).astype(np.uint8)
            if np.sum(binary_mask) > 0:
                masks_dict[class_id] = binary_mask
        encoded_pixels = combined_encode(masks_dict)
        results['ID'].append(image_id)
        results['encoded_pixels'].append(encoded_pixels)
submission_df = pd.DataFrame(results)
submission_df['encoded_pixels'] = submission_df['encoded_pixels'].fillna('')
submission_df.loc[submission_df['encoded_pixels'] == '', 'encoded_pixels'] = '<empty>'
submission_df = submission_df.sort_values('ID')
submission_df.to_csv(output_file, index=False)
print(f"Submission saved to {output_file}")
print(f"Total entries: {len(submission_df)}")
print(f"Empty predictions: {(submission_df['encoded_pixels'] == '<empty>').sum()}")
return submission_df

# Set device
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

# Load the trained model with correct encoder
model = get_enhanced_unet(encoder_name='resnet50', encoder_weights=None,
                           in_channels=1, out_channels=3)

# Load trained model weights (strict=False to handle key mismatches)
model.load_state_dict(torch.load('best_model.pth', map_location=device),
                      strict=False)

# Move model to device
model = model.to(device)
model.eval()

```

```
# Generate submission
generate_submission(
    model=model,
    test_dir='/content/content/kaggle_dataset_3/test',
    output_file='submission.csv',
    batch_size=16,
    device=device
)
```

Generating predictions...

100% | 10/10 [00:02<00:00, 3.67it/s]

Submission saved to submission.csv

Total entries: 156

Empty predictions: 17

[21]:

	ID	encoded_pixels
0	image_000	1:98131 3 98809 3 99487 3 100165 8 100843 8 10...
1	image_001	1:14854 39 15409 39 15895 17 15914 3 15921 6 1...
2	image_002	<empty>
3	image_003	1:46999 35 47688 35 48369 51 49058 51 49747 51...
4	image_004	1:151467 3 152239 3 153005 18 153777 18 154549...
..
151	image_151	1:94983 35 95739 35 96489 47 97245 47 97995 62...
152	image_152	1:34822 12 35174 24 35531 24 35883 32 36237 36...
153	image_153	1:145439 24 146202 24 146965 24 147719 51 1484...
154	image_154	2:148380 21 149163 21 149943 31 150726 31 1514...
155	image_155	1:56417 4 56977 4 57539 7 58099 7 58659 22 586...

[156 rows x 2 columns]